

Efficient SpMV on GPUs

Gaurish Telang

SUNY Stony Brook

December 4th, 2012

Table Of Contents

- 1 The CUDA model
- 2 The CUDA C/C++ API
- 3 Sparse Matrix Formats
- 4 Sparse-Matrix Vector Multiplication
- 5 Performance

The CUDA model

Host and Device

- **Host**



- **Device**



The CUDA model: Introduction

- An application consists of a *sequential* CPU program that may execute parallel programs known as **kernels** on a GPU.
- A **kernel** is a computation that is executed by a potentially large number (1000's) of parallel threads.

The CUDA model: Introduction

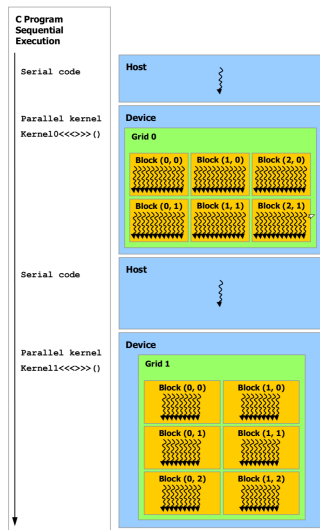
- Each thread runs the *same* scalar sequential program.
- Threads are organized into a grid of *thread-blocks*
- Threads of a **given** block can co-operate amongst themselves using
 - ▶ **Barrier** Synchronization.
 - ▶ **Atomic** operations
 - ▶ Per-block **shared** memory space *private* to each block.
- Each block can execute in *any order* relative to other blocks.
- Threads in different blocks *cannot* co-operate.

The CUDA model: Introduction

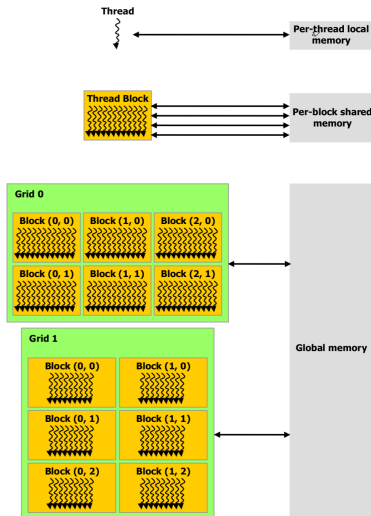
- Each thread runs the *same* scalar sequential program.
- Threads are organized into a grid of *thread-blocks*
- Threads of a **given** block can co-operate amongst themselves using
 - ▶ **Barrier** Synchronization.
 - ▶ **Atomic** operations
 - ▶ Per-block **shared** memory space *private* to each block.
- Each block can execute in *any order* relative to other blocks.
- Threads in different blocks *cannot* co-operate.

The CUDA model: Introduction

Program Execution




Memory Hierarchy



The CUDA model: Scalability

- Hardware is free to assign blocks to any processor at any time.
 - ▶ A kernel scales across any number of parallel processors.



The CUDA C/C++ API

Allocating and Releasing Device Memory

- **Memory Allocation**

```
int* ptr = 0 ;  
int  N   = 5000;  
int  num_bytes = N * sizeof( int );  
cudaMalloc( &ptr, num_bytes);
```

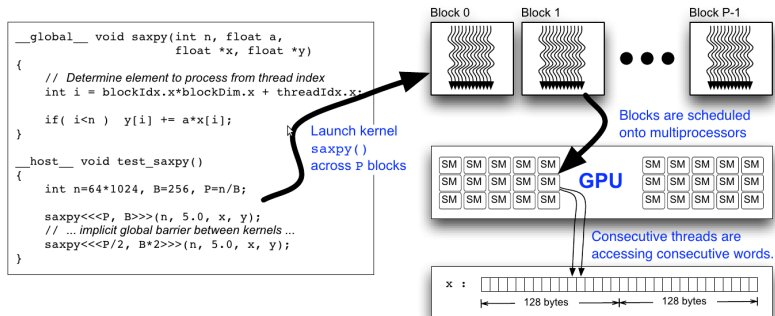
- **Memory transfer**

```
int* h = new int[5000]; //Fill the h array as needed  
cudaMemcpy(d , h , num_bytes , cudaMemcpyHostToDevice);  
// Do Computation on the device  
// ...  
cudaMemcpy(h , d , num_bytes , cudaMemcpyDeviceToHost);
```

- **Releasing Device Memory**

```
cudaFree(d);
```


Example CUDA Code: Simple SAXPY



Note: Calls to kernels are asynchronous by default.

The CUDA model: SIMT nature

- Thread creation, scheduling and management is performed entirely in hardware.
- To manage this large population of threads efficiently, the GPU employs a *SIMT* (Single Instruction Multiple Thread) architecture.
- Threads of a block are executed in *groups* of 32 called *warps*



The CUDA model: Warps

- A warp executes a *single* instruction at a time across **all** its threads.
- Threads of a warp are free to follow their own execution path and all such *execution divergence* is handled automatically in hardware.
- However it is **substantially** more efficient for threads to follow the *same* execution path for the bulk of the computation.


The CUDA model: Warps

- Threads of a warp are free to use *arbitrary* addresses when accessing off-chip memory with load/store operations.
- Accessing scattered locations results in *memory-divergence* and requires the processor to perform **one** memory transaction per-thread.

The CUDA model: Warps

- Non-contiguous access to memory is *detrimental* to memory bandwidth efficiency and to memory-bound computations.

```
--global--  
void saxpy_with_stride(int n, float a, float * x, float * y, int stride)  
{  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
  
    if (i < n)  
        y[i * stride] = a * x[i * stride] + y[i * stride];  
}
```



Relationship between stride and memory bandwidth in the `saxpy` and `daxpy` kernels.


Sparse Matrix Formats

Sparse Matrix Formats

- **Main Idea:** Encode the **non-zero** entries along with the **position** information.
- The encoding's *efficiency* depends on the *sparsity pattern*.
- Examples:
 - ▶ Diagonal (DIA)
 - ▶ ELLPACK (ELL)
 - ▶ Coordinate (COO)
 - ▶ Compressed Sparse Row (CSR)
 - ▶ Hybrid (ELL+COO)
 - ▶ Packet (PKT)
 - and others (including hybrids) ...

Diagonal Format: Introduction

- **Appropriate** when non-zero values are restricted to a **small** number of matrix diagonals,
- **Efficiently** encodes matrices arising from application of stencils to **regular** grids.



Diagonal Format: An Example

$$A = \begin{bmatrix} 1 & 7 & 0 & 0 \\ 0 & 2 & 8 & 0 \\ 5 & 0 & 3 & 9 \\ 0 & 6 & 0 & 4 \end{bmatrix}$$

$$\text{data} = \begin{bmatrix} * & 1 & 7 \\ * & 2 & 8 \\ 5 & 3 & 9 \\ 6 & 4 & * \end{bmatrix}$$

$$\text{offsets} = [-1 \quad 0 \quad 1]$$

Arrays **data** and **offsets** comprise the DIA representation of **A**.


Diagonal Format: + and -'s

Advantages:

- Row and column indices of each non-zero are defined *implicitly*. This *reduces* memory footprint and data-transfer during an SpMV operation.
- Good for stencils applied to *regular* grids.



Disadvantages:

- Allocates storage of values 'outside' the matrix and *explicitly* stores zero values along the diagonal.
- Not good for matrices with *arbitrary* sparsity patterns.



ELLPACK format: Introduction

- **Appropriate** for an M -by- N matrix with a *maximum* of K non-zeros per row where K is *small*.
- **Efficiently** encodes sparse matrices which have *roughly the same* number of non-zeros per row.



ELLPACK Format: An Example

$$A = \begin{bmatrix} 1 & 7 & 0 & 0 \\ 0 & 2 & 8 & 0 \\ 5 & 0 & 3 & 9 \\ 0 & 6 & 0 & 4 \end{bmatrix}$$

$$\text{data} = \begin{bmatrix} 1 & 7 & * \\ 2 & 8 & * \\ 5 & 3 & 9 \\ 6 & 4 & * \end{bmatrix}$$

$$\text{indices} = \begin{bmatrix} 0 & 1 & * \\ 1 & 2 & * \\ 0 & 2 & 3 \\ 6 & 4 & * \end{bmatrix}$$

Arrays `data` and `indices` comprise the ELL representation of `A`.


ELLPACK Format: + and -'s

Advantages:

- More general than DIA
- Good for matrices obtained from *semi-structured* and well-behaved *unstructured* meshes .

Disadvantages:

- For unstructured meshes, the *ratio* between **max** number of non-zeros per row and the **average** might be arbitrarily large.
- Hence a *vast* majority of the entries in the data and indices array will be *wasted*.



COO format

- Consists of **three** separate arrays:

- ▶ **row** row-indices
- ▶ **col** column-indices
- ▶ **data** non-zero values

$$A = \begin{bmatrix} 1 & 7 & 0 & 0 \\ 0 & 2 & 8 & 0 \\ 5 & 0 & 3 & 9 \\ 0 & 6 & 0 & 4 \end{bmatrix}$$

- Required storage *always proportional* to the number of non-zeros. unlike DIA and ELL

```
row = [0  0  1  1  2  2  2  3  3]
col = [0  1  1  2  0  2  3  1  3]
data = [1  7  2  8  5  3  9  6  4]
```

- Row-array is kept *sorted* ensuring entries with the same row-index are stored contiguously

CSR format

- A popular and *general purpose* sparse matrix representation
- Can be viewed as a *natural extension* of the (sorted) COO representation with a simple **compression** scheme applied to the (oft-repeated) row indices
- Consists of **three** separate arrays:
 - ▶ **ptr** new-row start pointers
 - ▶ **indices** column-indices
 - ▶ **data** non-zero values
- Required storage *always proportional* to the number of non-zeros. like the COO

$$A = \begin{bmatrix} 1 & 7 & 0 & 0 \\ 0 & 2 & 8 & 0 \\ 5 & 0 & 3 & 9 \\ 0 & 6 & 0 & 4 \end{bmatrix}$$

$\text{ptr} = [0 \quad 2 \quad 4 \quad 7 \quad 9]$
 $\text{indices} = [0 \quad 1 \quad 1 \quad 2 \quad 0 \quad 2 \quad 3 \quad 1 \quad 3]$
 $\text{data} = [1 \quad 7 \quad 2 \quad 8 \quad 5 \quad 3 \quad 9 \quad 6 \quad 4]$

CSR format

- A popular and *general purpose* sparse matrix representation
- Can be viewed as a *natural extension* of the (sorted) COO representation with a simple **compression** scheme applied to the (oft-repeated) row indices
- Consists of **three** separate arrays:
 - ▶ **ptr** new-row start pointers
 - ▶ **indices** column-indices
 - ▶ **data** non-zero values
- Required storage *always proportional* to the number of non-zeros. like the COO

$$A = \begin{bmatrix} 1 & 7 & 0 & 0 \\ 0 & 2 & 8 & 0 \\ 5 & 0 & 3 & 9 \\ 0 & 6 & 0 & 4 \end{bmatrix}$$

$\text{ptr} = [0 \quad 2 \quad 4 \quad 7 \quad 9]$
 $\text{indices} = [0 \quad 1 \quad 1 \quad 2 \quad 0 \quad 2 \quad 3 \quad 1 \quad 3]$
 $\text{data} = [1 \quad 7 \quad 2 \quad 8 \quad 5 \quad 3 \quad 9 \quad 6 \quad 4]$


HYB (= ELL + COO) format

- ELL format is suited for *vector machines*, but its efficiency rapidly degrades when the number of non-zeros per matrix varies a lot.
- *Storage efficiency* of COO is invariant to the distribution of non-zeros per row. This invariance **extends** to the cost of a SpMV product using COO.
- Split $A = B + C$ where
 - ▶ B contains the *typical* number of non-zeros per row and is stored in ELL.
 - ▶ C contains the remaining entries of the *exceptional* rows and is stored in COO


Packet(PKT) format

- Tailored for *symmetric* mesh-based matrices such as those in FEM.
- Idea
 - ▶ Reorder the non-zeros in the matrix such that non-zeros are 'concentrated' near the main diagonal.
 - ▶ Decompose this reordered matrix into a given number of fixed size partitions
 - ▶ Each partition is then stored in a packet data-structure. and each packet is assigned to a thread-block.

PKT format



A sample triangle mesh with 11 vertices partitioned into two sets.



Reordering rows and columns by partition concentrates nonzeros into diagonal submatrices.

PKT format

$$\text{pkt0} = \begin{bmatrix} (0,0) & (0,1) & (0,3) & (0,4) & * & * & * & * \\ (1,0) & (1,1) & (1,2) & (1,4) & (1,5) & * & * & * \\ (2,1) & (2,2) & (2,5) & (4,0) & (4,1) & (4,3) & (4,4) & (4,5) \\ (3,0) & (3,3) & (3,4) & (5,1) & (5,2) & (5,4) & (5,5) & * \end{bmatrix}$$
$$\text{pkt1} = \begin{bmatrix} (0,0) & (0,1) & (0,3) & (4,1) & (4,2) & (4,3) & (4,4) \\ (1,0) & (1,1) & (1,2) & (1,3) & (1,4) & * & * \\ (2,1) & (2,2) & (2,4) & * & * & * & * \\ (3,0) & (3,1) & (3,3) & (3,4) & * & * & * \end{bmatrix}$$

Packet storage for the two diagonal submatrices in Figure 13. The first submatrix has base index (0,0) while the second submatrix begins at (6,6). Rows of the `pkt0` and `pkt1` correspond to individual thread lanes.

Sparse-Matrix Vector Multiplication

Sparse Matrix Vector multiplication

- We consider the operation $y \leftarrow Ax + y$
- **Central idea**
Minimize the number of *independent* paths of execution and use the *coalescing* of threads whenever possible.

SpMV for the DIA format

```
__global__ void
spmv_dia_kernel(const int num_rows,
                const int num_cols,
                const int num_diags,
                const int * offsets,
                const float * data,
                const float * x,
                float * y)
{
    int row = blockDim.x * blockIdx.x + threadIdx.x;

    if(row < num_rows){
        float dot = 0;

        for(int n = 0; n < num_diags; n++){
            int col = row + offsets[n];
            float val = data[num_rows * n + row];

            if(col >= 0 && col < num_cols)
                dot += val * x[col];
        }

        y[row] += dot;
    }
}
```

```
data    [*  *  5  6  1  2  3  4  7  8  9  *]
```

Iteration 0	2	3
-------------	---	---

Iteration 1	0	1	2	3
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0

Iteration 2	0	1	2
-------------	---	---	---

SpMV for the ELL format

```
__global__ void
spmv_ell_kernel(const int num_rows,
                const int num_cols,
                const int num_cols_per_row,
                const int * indices,
                const float * data,
                const float * x,
                float * y)
{
    int row = blockDim.x * blockIdx.x + threadIdx.x;

    if(row < num_rows){
        float dot = 0;

        for(int n = 0; n < num_cols_per_row; n++){
            int col = indices[num_rows * n + row];
            float val = data[num_rows * n + row];

            if(val != 0)
                dot += val * x[col];
        }

        y[row] += dot;
    }
}
```

data	[1 2 5 6 7 8 3 4 * * 9 *]
indices	[0 1 0 6 1 2 2 4 * * 3 *]

Iteration 0	[0 1 2 3]	
Iteration 1	[0 1 2 3]
Iteration 2	[0 1 2 3]

SpMV for the CSR format(first-version)

```
__host__ void
spmv_csr_serial(const int num_rows,
                const int * ptr,
                const int * indices,
                const float * data,
                const float * x,
                float * y)
{
    for(int row = 0; row < num_rows; row++){
        float dot = 0;

        int row_start = ptr[row];
        int row_end   = ptr[row+1];

        for (int jj = row_start; jj < row_end; jj++)
            dot += data[jj] * x[indices[jj]];

        y[row] += dot;
    }
}
```

indices	[0 1 1 2 0 2 3 1 3]
data	[1 7 2 8 5 3 9 6 4]

Iteration 0	[0 1 2 3]
Iteration 1	[0 2 3]
Iteration 2	[2]

SpMV for the CSR format(second-version)

```

__global__ void
spmv_csr_vector_kernel(const int num_rows,
                       const int * ptr,
                       const int * indices,
                       const float * data,
                       const float * x,
                       float * y)
{
    __shared__ float vals[];

    int thread_id = blockDim.x * blockIdx.x + threadIdx.x; // global thread index
    int warp_id   = thread_id / 32;                        // global warp index
    int lane      = thread_id & (32 - 1);                  // thread index within the warp

    // one warp per row
    int row = warp_id;

    if (row < num_rows){
        int row_start = ptr[row];
        int row_end   = ptr[row+1];

        // compute running sum per thread
        vals[threadIdx.x] = 0;
        for(int jj = row_start + lane; jj < row_end; jj += 32)
            vals[threadIdx.x] += data[jj] * x[indices[jj]];

        // parallel reduction in shared memory
        if (lane < 16) vals[threadIdx.x] += vals[threadIdx.x + 16];
        if (lane < 8)  vals[threadIdx.x] += vals[threadIdx.x + 8];
        if (lane < 4)  vals[threadIdx.x] += vals[threadIdx.x + 4];
        if (lane < 2)  vals[threadIdx.x] += vals[threadIdx.x + 2];
        if (lane < 1)  vals[threadIdx.x] += vals[threadIdx.x + 1];

        // first thread writes the result
        if (lane == 0)
            y[row] += vals[threadIdx.x];
    }
}

```

	indices	[0 1 1 2 0 2 3 1 3]
	data	[1 7 2 8 5 3 9 6 4]
Warp 0		[0 0]
Warp 1		[1 1]
Warp 2		[2 2 2]
Warp 3		[3 3]

SpMV for the PKT format

```
__device__ void
process_packet(const int work_per_thread,
               const int * index_packet,
               float * data_packet,
               const float * x_local,
               float * y_local)
{
    for(int i = 0; i < work_per_thread; i++){
        // offset into packet arrays
        int pos = blockDim.x * i + threadIdx.x;

        // row and column indices are packed in one 32-bit word
        int packed_index = index_array[pos];

        // row and column stored in upper and lower half-words
        int row = packed_index >> 16;
        int col = packed_index & 0xFFFF;

        float val = data_array[pos];

        y_local[row] += val * x_local[col];
    }
}
```

A block of threads computes the matrix-vector product for a single packet. Here `y_local` reside in shared memory for low-latency access.

Performance

Performance

- Different classes of matrices highlight the strength and weaknesses of the sparse formats and their associated SpMV algorithms.
- Performance data are measured in terms of
 - ▶ Speed of execution (GFLOP/s)
 - ▶ Memory bandwidth utilization (GBytes/s)
- Measurements do *not* include time spent *transferring* data between the CPU and the GPU since we are trying to measure just the SpMV performance.
- The reported figures represent an *average* (arithmetic-mean) of **500** SpMV operations.


Structured Matrices

Performance: Structured Matrices


Matrix	Grid	Diagonals	Rows	Columns	Nonzeros
Laplace 3pt	$(1,000,000)$	3	1,000,000	1,000,000	2,999,998
Laplace 5pt	$(1,000)^2$	5	1,000,000	1,000,000	4,996,000
Laplace 7pt	$(100)^3$	7	1,000,000	1,000,000	6,940,000
Laplace 9pt	$(1,000)^2$	9	1,000,000	1,000,000	8,988,004
Laplace 27pt	$(100)^3$	27	1,000,000	1,000,000	26,463,592

Structured matrices used for performance testing.

Performance: Structured Matrices: Execution Speed










Performance: Structured Matrices: Memory Bandwidth










Unstructured Matrices

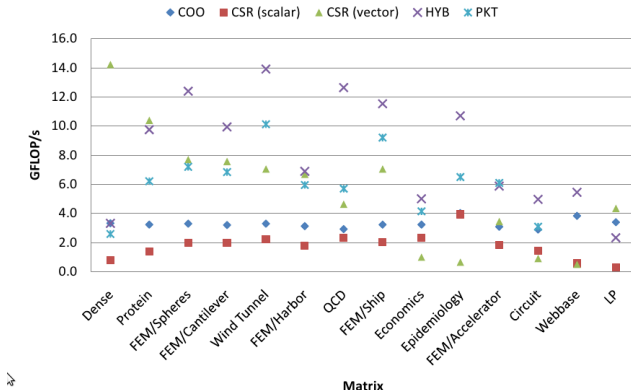
Performance: Unstructured Matrices: DataSet

spyplot	Name	Dimensions	Nonzeros (nnz/row)	Description
	Dense	2K x 2K	4.0M (2K)	Dense matrix in sparse format
	Protein	36K x 36K	4.3M (119)	Protein data bank 1HYS
	FEM / Spheres	83K x 83K	6.0M (72)	FEM concentric spheres
	FEM / Cantilever	62K x 62K	4.0M (65)	FEM cantilever
	Wind Tunnel	218K x 218K	11.6M (53)	Pressurized wind tunnel
	FEM / Harbor	47K x 47K	2.37M (50)	3D CFD of Charleston harbor
	QCD	49K x 49K	1.90M (39)	Quark propagators (QCD/LGT)


Performance: Unstructured Matrices: DataSet (ctd...)

	FEM/Ship	141K x 141K	3.98M (28)	FEM Ship section/detail
	Economics	207K x 207K	1.27M (6)	Macroeconomic model
	Epidemiology	526K x 526K	2.1M (4)	2D Markov model of epidemic
	FEM / Accelerator	121K x 121K	2.62M (22)	Accelerator cavity design
	Circuit	171K x 171K	959K (6)	Motorola circuit simulation
	webbase	1M x 1M	3.1M (3)	Web connectivity matrix
	LP	4K x 1.1M	11.3M (2825)	Railways set cover Constraint matrix

Performance: Unstructured Matrices: Speed Of Execution



Performance: Unstructured Matrices: Memory Bandwidth

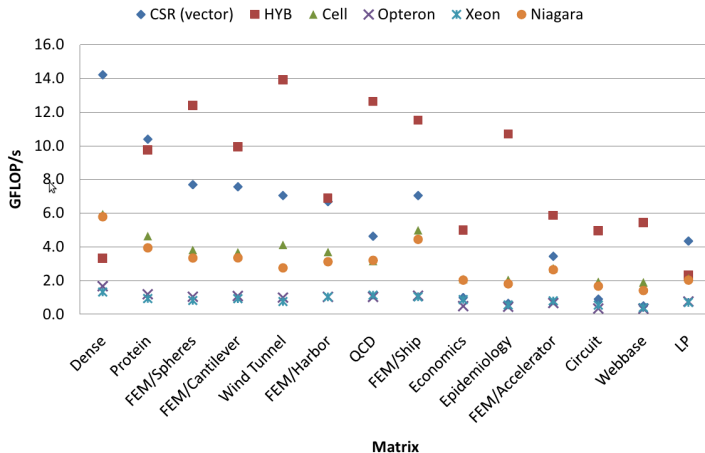


Performance: Comparisons With Multicore


Platform Specifications

Name	Sockets	Cores	Clock (GHz)	Description
Cell	1	8 (SPEs)	3.2	IBM QS20 Cell Blade (half)
Opteron	1	2	2.2	AMD Opteron 2214
Xeon	1	4	2.3	Intel Clovertown
Niagara	1	8	1.4	Sun Niagara2
Dual Cell	2	16 (SPEs)	3.2	IBM QS20 Cell Blade (full)
Dual Opteron	2	4	2.2	2 x AMD Opteron 2214
Dual Xeon	2	8	2.3	2 x Intel Clovertown

Performance: Comparison With Single Socket



Performance: Comparison With Double Socket



Questions?